

МОСКОВСКИЙ ГОСУДАРСТВЕННЫЙ ТЕХНИЧЕСКИЙ УНИВЕРСИТЕТ  
им. Н.Э. Баумана

Кафедра «Систем обработки информации и управления»

ОТЧЕТ

**Лабораторная работа №6**  
по курсу «Методы машинного обучения»

Тема: «Ансамбли моделей машинного обучения»

ИСПОЛНИТЕЛЬ:

группа ИУ5-22М

\_\_Бабин В.Е.\_\_\_\_\_  
ФИО

\_\_\_\_\_  
подпись

"\_\_"\_\_\_\_2020 г.

ПРЕПОДАВАТЕЛЬ:

\_\_\_\_\_  
ФИО

\_\_\_\_\_  
подпись

"\_\_"\_\_\_\_2020 г.

Москва - 2020

---

## Задание:

1. Выберите набор данных (датасет) для решения задачи классификации или регрессии.
2. В случае необходимости проведите удаление или заполнение пропусков и кодирование категориальных признаков.
3. С использованием метода `train_test_split` разделите выборку на обучающую и тестовую.
4. Обучите две ансамблевые модели. Оцените качество моделей с помощью одной из подходящих для задачи метрик. Сравните качество полученных моделей.
5. Произведите для каждой модели подбор значений одного гиперпараметра. В зависимости от используемой библиотеки можно применять функцию `GridSearchCV`, использовать перебор параметров в цикле, или использовать другие методы.
6. Повторите пункт 4 для найденных оптимальных значений гиперпараметров. Сравните качество полученных моделей с качеством моделей, полученных в пункте 4.

# lab6

May 19, 2020

```
[1]: import pandas as pd
import numpy as np
import io
import requests
from sklearn.impute import SimpleImputer

# mushrooms dataset
url = "https://www.wolframcloud.com/obj/d0c0084e-0b60-46db-a4d9-beb33412905e"
s = requests.get(url).content
data = pd.read_csv(io.StringIO(s.decode('utf-8')))
data.head()
```

```
[1]:   CapShape CapSurface CapColor ... Population Habitat      Class
0   convex      smooth   brown ...   scattered   urban  poisonous
1   convex      smooth  yellow ...   numerous  grasses    edible
2    bell      smooth   white ...   numerous  meadows    edible
3   convex      scaly   white ...   scattered   urban  poisonous
4   convex      smooth   gray ...   abundant  grasses    edible
```

[5 rows x 23 columns]

## Dataset preparation

```
[2]: rows, columns = data.shape
print('rows = {}; columns = {}'.format(rows, columns), '\n\n')

# go through each dataset column and check unique values to find empty one
→ (like NaN or Missing[])
for col in data.columns:
    print('{}: {}'.format(col, data[col].unique()))
```

rows = 8124; columns = 23

CapShape: ['convex' 'bell' 'sunken' 'flat' 'knobbed' 'conical']  
CapSurface: ['smooth' 'scaly' 'fibrous' 'grooves']

```

CapColor: ['brown' 'yellow' 'white' 'gray' 'red' 'pink' 'buff' 'purple'
'cinnamon'
'green']
Bruises: [ True False]
Odor: ['pungent' 'almond' 'anise' 'none' 'foul' 'creosote' 'fishy' 'spicy'
'musty']
GillAttachment: ['free' 'attached']
GillSpacing: ['close' 'crowded']
GillSize: ['narrow' 'broad']
GillColor: ['black' 'brown' 'gray' 'pink' 'white' 'chocolate' 'purple' 'red'
'buff'
'green' 'yellow' 'orange']
StalkShape: ['enlarging' 'tapering']
StalkRoot: ['equal' 'club' 'bulbous' 'rooted' 'Missing[]']
StalkSurfaceAboveRing: ['smooth' 'fibrous' 'silky' 'scaly']
StalkSurfaceBelowRing: ['smooth' 'fibrous' 'scaly' 'silky']
StalkColorAboveRing: ['white' 'gray' 'pink' 'brown' 'buff' 'red' 'orange'
'cinnamon' 'yellow']
StalkColorBelowRing: ['white' 'pink' 'gray' 'buff' 'brown' 'red' 'yellow'
'orange' 'cinnamon']
VeilType: ['partial']
VeilColor: ['white' 'brown' 'orange' 'yellow']
RingNumber: [1 2 0]
RingType: ['pendant' 'evanescent' 'large' 'flaring' 'none']
SporePrintColor: ['black' 'brown' 'purple' 'chocolate' 'white' 'green' 'orange'
'yellow'
'buff']
Population: ['scattered' 'numerous' 'abundant' 'several' 'solitary' 'clustered']
Habitat: ['urban' 'grasses' 'meadows' 'woods' 'paths' 'waste' 'leaves']
Class: ['poisonous' 'edible']

```

```

[3]: """
on initial viewing it seems that we have single column
with absence of values in rows: this column is StalkRoot and
absence of values is indicated like Missing[]
"""

# Take a look at columns more precisely to ensure that this column is single_
→with absence of values
for col in data.columns:
    # Missing[] amount
    temp_null_count = data[data[col] == 'Missing[]'].shape[0]
    dt = str(data[col].dtype)
    if temp_null_count>0 and (dt=='object'):
        temp_perc = round((temp_null_count / rows) * 100.0, 2)
        print('Column {}. Data type {}. amount of Missing[] values {}, {}%.'.
→format(col, dt, temp_null_count, temp_perc))

```

Column StalkRoot. Data type object. amount of Missing[] values 2480, 30.53%.  
/usr/local/lib/python3.6/dist-packages/pandas/core/ops/array\_ops.py:253:  
FutureWarning: elementwise comparison failed; returning scalar instead, but in  
the future will perform elementwise comparison  
res\_values = method(rvalues)

```
[4]: data['StalkRoot'].unique()
```

```
[4]: array(['equal', 'club', 'bulbous', 'rooted', 'Missing[]'], dtype=object)
```

```
[5]: # impute data with most frequent values  
imputation = SimpleImputer(missing_values='Missing[]', strategy='most_frequent')  
data_imputed = imputation.fit_transform(data[['StalkRoot']])  
np.unique(data_imputed)
```

```
[5]: array(['bulbous', 'club', 'equal', 'rooted'], dtype=object)
```

```
[6]: # put imputed data in our dataset  
for i in range(rows):  
    data['StalkRoot'][i] = data_imputed[i][0]  
data['StalkRoot'].unique()
```

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:3:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
This is separate from the ipykernel package so we can avoid doing imports  
until

```
[6]: array(['equal', 'club', 'bulbous', 'rooted'], dtype=object)
```

```
[7]: # now we'll transform categorical columns to columns with values in [0, 1]  
object_columns = []  
  
# initially, find amount of unique values for each column  
# and add categorical columns to object_columns  
for column in data.columns:  
    dt = str(data[column].dtype)  
    amount_unique = len(pd.unique(data[column]))  
    print('{}: {}, {}'.format(column, amount_unique, dt))  
    if dt == 'object':  
        object_columns.append({'column': column, 'amount': amount_unique})
```

CapShape: 6, object  
CapSurface: 4, object

```

CapColor: 10, object
Bruises: 2, bool
Odor: 9, object
GillAttachment: 2, object
GillSpacing: 2, object
GillSize: 2, object
GillColor: 12, object
StalkShape: 2, object
StalkRoot: 4, object
StalkSurfaceAboveRing: 4, object
StalkSurfaceBelowRing: 4, object
StalkColorAboveRing: 9, object
StalkColorBelowRing: 9, object
VeilType: 1, object
VeilColor: 4, object
RingNumber: 3, int64
RingType: 5, object
SporePrintColor: 9, object
Population: 6, object
Habitat: 7, object
Class: 2, object

```

```

[0]: from sklearn.preprocessing import OneHotEncoder
     ohe = OneHotEncoder()

```

```

[0]: """
     after oneHot encoding for single value we'll have something like this [0.0, 1.
     ↪0, 0.0, 0.0, ...]
     but we need to have just a number - so this function will normalize it with
     ↪next formula:
     norm_val = (N - i) / N, where
     N - amount of unique values for our column
     i - index of 1.0 in values [0.0, 1.0, 0.0, 0.0, ...] before normalizing
     so for example if we'll have 10 different values - then in normalized view it
     ↪will variating
     from 0.1 (if index = 9) to 1.0 (if index = 0)
     """
     def from_bytes_to_num(col_in_arr, unique_amount):
         normalized_col = []
         for value in col_in_arr:
             normalized_value = (unique_amount - np.where(value == 1.0)[0][0]) /
             ↪unique_amount
             normalized_col.append(float("{0:.4f}".format(normalized_value))) #
             ↪digits after float point
         return normalized_col

```

```

normalized_data = []
for col in object_columns:
    unique_amount = col['amount']
    col_name = col['column']

    #encode with oneHot
    column_after_encoding = ohe.fit_transform(data[[col_name]])

    #fetch it to array
    col_in_arr = column_after_encoding.toarray()

    #normalizing column values
    normalized_col = from_bytes_to_num(col_in_arr, unique_amount)

    normalized_data.append({'column': col_name, 'data': normalized_col})

```

```

[10]: #set normalized values for general dataset
for col in normalized_data:
    col_in_dataFrame = pd.DataFrame(data={col['column']: col['data']})
    data[col['column']] = col_in_dataFrame

data

```

```

[10]:
   CapShape  CapSurface  CapColor  ...  Population  Habitat  Class
0      0.6667         0.25        1.0  ...      0.5000     0.4286    0.5
1      0.6667         0.25        0.1  ...      0.6667     1.0000    1.0
2      1.0000         0.25        0.2  ...      0.6667     0.7143    1.0
3      0.6667         0.50        0.2  ...      0.5000     0.4286    0.5
4      0.6667         0.25        0.7  ...      1.0000     1.0000    1.0
...      ...          ...      ...  ...      ...      ...      ...
8119    0.3333         0.25        1.0  ...      0.8333     0.8571    1.0
8120    0.6667         0.25        1.0  ...      0.3333     0.8571    1.0
8121    0.5000         0.25        1.0  ...      0.8333     0.8571    1.0
8122    0.3333         0.50        1.0  ...      0.3333     0.8571    0.5
8123    0.6667         0.25        1.0  ...      0.8333     0.8571    1.0

```

[8124 rows x 23 columns]

## Data splitting

```

[11]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_score, train_test_split

dcopy = data.copy()

X = dcopy.drop("Class", axis=1)

```

```
y = dcopy["Class"]
print(X.head(), "\n")
print(y.head())
```

	CapShape	CapSurface	CapColor	...	SporePrintColor	Population	Habitat
0	0.6667	0.25	1.0	...	1.0000	0.5000	0.4286
1	0.6667	0.25	0.1	...	0.8889	0.6667	1.0000
2	1.0000	0.25	0.2	...	0.8889	0.6667	0.7143
3	0.6667	0.50	0.2	...	1.0000	0.5000	0.4286
4	0.6667	0.25	0.7	...	0.8889	1.0000	1.0000

[5 rows x 22 columns]

```
0    0.5
1    1.0
2    1.0
3    0.5
4    1.0
```

Name: Class, dtype: float64

```
[12]: # divide train and test selections
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.25, random_state=1)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(6093, 22)
(2031, 22)
(6093,)
(2031,)
```

## Model training

```
[0]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import median_absolute_error, r2_score

def test_model(model):
    print("mean_absolute_error:",
          mean_absolute_error(y_test, model.predict(X_test)))
    print("median_absolute_error:",
          median_absolute_error(y_test, model.predict(X_test)))
    print("r2_score:",
          r2_score(y_test, model.predict(X_test)))
```



## Random forest

```
[29]: from sklearn.ensemble import RandomForestRegressor
```

```
ran_80 = RandomForestRegressor(n_estimators=80)
ran_80.fit(X_train, y_train)
```

```
[29]: RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             max_samples=None, min_impurity_decrease=0.0,
                             min_impurity_split=None, min_samples_leaf=1,
                             min_samples_split=2, min_weight_fraction_leaf=0.0,
                             n_estimators=80, n_jobs=None, oob_score=False,
                             random_state=None, verbose=0, warm_start=False)
```

```
[30]: test_model(ran_80)
```

```
mean_absolute_error: 0.0
median_absolute_error: 0.0
r2_score: 1.0
```

## Gradient Boost

```
[31]: from sklearn.ensemble import GradientBoostingRegressor
```

```
gr_80 = GradientBoostingRegressor(n_estimators=80)
gr_80.fit(X_train, y_train)
```

```
[31]: GradientBoostingRegressor(alpha=0.9, ccp_alpha=0.0, criterion='friedman_mse',
                                 init=None, learning_rate=0.1, loss='ls', max_depth=3,
                                 max_features=None, max_leaf_nodes=None,
                                 min_impurity_decrease=0.0, min_impurity_split=None,
                                 min_samples_leaf=1, min_samples_split=2,
                                 min_weight_fraction_leaf=0.0, n_estimators=80,
                                 n_iter_no_change=None, presort='deprecated',
                                 random_state=None, subsample=1.0, tol=0.0001,
                                 validation_fraction=0.1, verbose=0, warm_start=False)
```

```
[32]: test_model(gr_80)
```

```
mean_absolute_error: 0.006802494393104271
median_absolute_error: 0.0010510475336632519
r2_score: 0.9900409523010036
```

n

## Hyperparameter n selection

## Random forest

```
[46]: param_range = np.arange(50, 170, 10)
      tuned_parameters = [{'n_estimators': param_range}]
      tuned_parameters
```

```
[46]: [{'n_estimators': array([ 50,  60,  70,  80,  90, 100, 110, 120, 130, 140, 150,
 160])}]
```

```
[47]: from sklearn.model_selection import GridSearchCV
      from sklearn.model_selection import ShuffleSplit


      gs = GridSearchCV(RandomForestRegressor(), tuned_parameters,
                        cv=ShuffleSplit(n_splits=10), scoring="r2",
                        return_train_score=True, n_jobs=-1)
      gs.fit(X, y)
      gs.best_estimator_
```

```
[47]: RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                           max_depth=None, max_features='auto', max_leaf_nodes=None,
                           max_samples=None, min_impurity_decrease=0.0,
                           min_impurity_split=None, min_samples_leaf=1,
                           min_samples_split=2, min_weight_fraction_leaf=0.0,
                           n_estimators=80, n_jobs=None, oob_score=False,
                           random_state=None, verbose=0, warm_start=False)
```

```
[48]: import matplotlib.pyplot as plt
      %matplotlib inline


      plt.plot(param_range, gs.cv_results_["mean_train_score"]);
```

```
[48]:
```



```
[49]: plt.plot(param_range, gs.cv_results_["mean_test_score"]);
```

[49]:



```
[50]: reg = gs.best_estimator_  
      reg.fit(X_train, y_train)  
      test_model(reg)
```

```
mean_absolute_error: 6.154603643525363e-05  
median_absolute_error: 0.0  
r2_score: 0.9999753811019957
```


### Gradient Boost

```
[28]: gs = GridSearchCV(GradientBoostingRegressor(), tuned_parameters,  
                        cv=ShuffleSplit(n_splits=10), scoring="r2",  
                        return_train_score=True, n_jobs=-1)  
  
gs.fit(X, y)  
gs.best_estimator_
```

```
[28]: GradientBoostingRegressor(alpha=0.9, ccp_alpha=0.0, criterion='friedman_mse',  
                                init=None, learning_rate=0.1, loss='ls', max_depth=3,  
                                max_features=None, max_leaf_nodes=None,  
                                min_impurity_decrease=0.0, min_impurity_split=None,  
                                min_samples_leaf=1, min_samples_split=2,  
                                min_weight_fraction_leaf=0.0, n_estimators=200,  
                                n_iter_no_change=None, presort='deprecated',  
                                random_state=None, subsample=1.0, tol=0.0001,  
                                validation_fraction=0.1, verbose=0, warm_start=False)
```


```
[33]: plt.plot(param_range, gs.cv_results_["mean_train_score"]);
```

[33]:



```
[34]: plt.plot(param_range, gs.cv_results_["mean_test_score"]);
```

[34]:



```
[35]: reg = gs.best_estimator_  
reg.fit(X_train, y_train)  
test_model(reg)
```

```
mean_absolute_error: 0.0028189626630988947  
median_absolute_error: 0.0005545827653701263  
r2_score: 0.9980449702667502
```

[0]: